

Intelligent Multi-Agent Coordination with Adaptive Knowledge Structures for Complex Web Workflows

Zubair Ahmed Rafi

SUST Research Center

zubairahmedrafi37@gmail.com

Md. Rafi Alam

Rajshahi University of Engineering & Technology

rafialam6610@gmail.com

Md. Ibrahim Ibna Khalil

Daffodil International University (DIU)

khalil22205101816@diu.edu.bd

Abstract—Autonomous web automation frequently fails when single agents attempt complex, long-horizon workflows. These failures due to sequential bottlenecks and a lack of structured memory, rendering agents unable to learn from past mistakes. To address this, we propose a multi-agent orchestration framework that replaces static prompt chains with an Adaptive Knowledge Graph. By modeling task decomposition as a Markov Decision Process (MDP) driven by Bayesian inference, our system dynamically updates dependency probabilities based on execution history. This allows the orchestrator to “learn” the optimal decomposition path and agent assignment over time. Crucially, we enforce a fixed iteration bound to provide strict termination guarantees, solving the infinite-loop problem common in recursive agent architectures. Evaluation on multi-modal travel planning tasks demonstrates that our approach is not only more robust but increasingly efficient. This work bridges the gap between probabilistic reasoning and large language model (LLM) agents, offering a scalable architecture for self-optimizing web automation.

Index Terms—multi-agent systems, knowledge graphs, Markov decision processes, Bayesian inference, task orchestration, web automation

I. INTRODUCTION

The transition from passive information retrieval to autonomous task execution represents a fundamental shift in web-based artificial intelligence. While Large Language Models (LLMs) demonstrate remarkable capability in generating code and natural language, their application to complex, long-horizon web automation remains brittle. A composite objective, such as coordinating a multi-modal travel itinerary, is not merely a sequence of API calls but a dynamic workflow characterized by stochastic dependencies, temporal coupling, and non-deterministic service availability.

Current approaches largely fall into two paradigms: monolithic single-agent systems, which suffer from context window exhaustion and sequential bottlenecks, and uncoordinated multi-agent frameworks that distribute tasks but lack shared memory or structural coherence [1]. A critical limitation in both approaches is the absence of probabilistic state tracking. When a subtask fails or returns suboptimal data, stateless agents typically lack the historical context to adjust their decomposition strategy, resulting in redundant retries or catastrophic propagation of errors.

In this paper, we propose a structured orchestration framework that bridges the gap between neural reasoning and symbolic planning. We argue that robust web automation

requires three architectural primitives: (1) a persistent memory structure that encodes task dependencies, (2) a probabilistic mechanism to model uncertainty in subtask execution, and (3) an adaptive resource allocation strategy.

To realize this, we introduce a system driven by an Adaptive Knowledge Graph (AKG), which serves as a dynamic schema of the task space. Unlike static planning graphs, our AKG evolves via Bayesian updates, allowing the system to learn optimal decomposition paths from execution history. We formalize the orchestration problem as a Markov Decision Process (MDP), where agent selection is optimized through a scoring function that weighs specialization against current workload. This probabilistic approach allows for mathematically bounded error recovery, ensuring that the system converges toward successful workflows while respecting strict termination guarantees.

II. RELATED WORK

A. Multi-Agent Task Decomposition and Coordination

Task decomposition in multi-agent systems has been explored through symbolic, hierarchical, and role-based formalisms. Shah et al. [2] (Liu et al.) proposed a framework for learning symbolic task decompositions using PDDL solvers to empower LLMs. Ardon et al. [3] (Toro Icarte et al.) utilized reward machines to provide a formal decomposition of cooperative tasks, guiding agents through stateful automata. Other works explore dynamic role discovery through Bayesian policy-search methods, such as the approach of Wilson and Fern [4], which infers latent roles from experience and enables coordinated behavior without predefined role assignments. The classical Contract Net Protocol [5] remains foundational for distributed task allocation, offering early insights into agent bidding mechanisms.

B. Reinforcement Learning with Task Decomposition

Reinforcement learning approaches have integrated hierarchical decomposition to reduce complexity in cooperative settings. Sun et al. [6] (Wang et al.) proposed *Voyager*, an embodied agent that uses an automatic curriculum to decompose tasks into manageable sub-routines. In robotic manipulation, hierarchical VLA models like RT-2 [7] have been applied to transfer web knowledge into independent control primitives. While effective in physical control domains, these methods often depend on massive pre-training or manually

defined hierarchies, limiting their flexibility in open-ended web environments.

C. Knowledge Graphs and Structured Memory

Knowledge graphs (KGs) have gained traction as structured memory representations that offer distinct advantages over unstructured vector retrieval (RAG). Qin and Lu [8] (Pan et al.) provided a roadmap for unifying LLMs with KGs to improve planning capabilities. In reasoning architectures, Wu et al. [9] (Besta et al.) introduced *Graph of Thoughts*, modeling dependencies as a directed graph to optimize prompt execution. Recently, Yang et al. [10] (Wang et al.) introduced *KnowAgent*, an adaptive hierarchical KG for multi-agent collaboration. Unlike vector stores, KGs explicitly model the *relational* dependencies between subtasks, which is critical for maintaining state in long-horizon workflows.

D. LLM Agents for Web Automation

The application of Large Language Models to web automation has evolved from simple prompt-chaining to autonomous agentic workflows. Foundational reasoning frameworks like *Chain-of-Thought* [11] and *ReAct* [12] established the paradigm of interleaving reasoning with API execution. However, benchmarks such as WebArena [13] highlighted the fragility of these single-agent loops when facing complex, multi-step objectives due to context window exhaustion. While architectures like *Generative Agents* [14] introduced memory reflection mechanisms for social simulation, they lack the rigid termination guarantees required for functional web tasks. Our work addresses this by combining the structured memory of KGs with the rigorous termination bounds of formal orchestration.

E. Limitations in Prior Work

Despite progress, existing works exhibit several gaps: (1) many rely on static decomposition structures that fail to adapt to dynamic web changes; (2) uncertainty in decomposition success is seldom modeled explicitly; (3) agent–task matching mechanisms are typically heuristic rather than formally scored; and (4) termination guarantees are rarely addressed, resulting in infinite retry loops. Our framework addresses these gaps by introducing probabilistic edge-updating, explicit scoring for agent–task compatibility, and controlled termination bounds.

III. PROBLEM FORMULATION

A. Task Model

Let $\mathcal{T} = \{t_1, t_2, \dots, t_n\}$ represent the space of tasks. A complex task $T \in \mathcal{T}$ can be decomposed into a sequence of subtasks:

$$T = \langle s_1, s_2, \dots, s_k \rangle$$

where each s_i is an atomic subtask. Subtasks may have dependencies represented as a directed acyclic graph (DAG):

$$\mathcal{D} = (\mathcal{S}, \mathcal{E})$$

where \mathcal{S} is the set of subtasks and $\mathcal{E} \subseteq \mathcal{S} \times \mathcal{S}$ represents dependency edges. An edge $(s_i, s_j) \in \mathcal{E}$ indicates that s_j depends on the output of s_i .

B. Agent Model

The system maintains a pool of agents $\mathcal{A} = \{a_1, a_2, \dots, a_m\}$, where each agent a_i has:

- A capability vector $\mathbf{c}_i \in \mathbb{R}^d$ representing skills
- A current workload $w_i \in [0, 1]$
- A performance history \mathcal{H}_i containing past task executions

C. Knowledge Graph Structure

Our adaptive knowledge graph $\mathcal{G} = (\mathcal{N}, \mathcal{R}, \mathbf{W})$ consists of:

- \mathcal{N} : Node set representing atomic subtasks and task types
- \mathcal{R} : Edge set representing relationships (decomposition, dependency, sequence)
- \mathbf{W} : Weight matrix where w_{ij} represents the learned probability/utility of edge (n_i, n_j)

The graph maintains three types of edges:

- 1) *Decomposition edges*: Connect composite tasks to their subtasks
- 2) *Dependency edges*: Represent data or control flow between subtasks
- 3) *Sequence edges*: Indicate temporal ordering with execution probabilities

D. Objective

Given a complex task T , our system aims to:

$$\begin{aligned} &\text{maximize} && P(\text{success} \mid T, \mathcal{A}, \mathcal{G}) \\ &\text{subject to} && \text{iterations}(s_i) \leq \kappa_{\max}, \forall s_i \\ & && \sum_j \text{load}(a_j) \leq |\mathcal{A}| \end{aligned}$$

where κ_{\max} is the maximum iteration bound and load constraints prevent agent oversubscription.

IV. METHODOLOGY

A. Probabilistic Task Decomposition

1) *Markov Decision Process Formulation*: We model task decomposition as an MDP $\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$ where:

- \mathcal{S} : State space representing partial task completion states. Each state $s \in \mathcal{S}$ encodes: (1) set of completed subtasks, (2) available next subtasks, (3) accumulated context
- \mathcal{A} : Action space representing subtask assignment decisions
- $\mathcal{P} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$: Transition probability function
- $\mathcal{R} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$: Reward function based on subtask completion
- $\gamma \in [0, 1]$: Discount factor

The transition probability $P(s' \mid s, a)$ represents the likelihood of reaching state s' after taking action a in state s . These probabilities are initially set based on domain knowledge and updated through Bayesian inference as described below.

2) *Bayesian Probability Estimation*: For each edge (n_i, n_j) in the knowledge graph, we maintain a probability distribution over success rates. We use a Beta distribution as the prior:

$$P(\theta_{ij}) = \text{Beta}(\alpha_{ij}, \beta_{ij})$$

where θ_{ij} is the true success probability for transitioning from subtask n_i to n_j , and α_{ij}, β_{ij} are hyperparameters.

After observing k successes in n attempts, we update via Bayesian inference:

$$P(\theta_{ij} | \text{data}) = \text{Beta}(\alpha_{ij} + k, \beta_{ij} + n - k)$$

The expected success probability becomes:

$$\mathbb{E}[\theta_{ij}] = \frac{\alpha_{ij} + k}{\alpha_{ij} + \beta_{ij} + n}$$

This approach naturally handles the exploration–exploitation tradeoff through the posterior distribution’s variance.

3) *Decision Strategy*: At each decomposition step, the orchestrator selects the next subtask assignment using a policy $\pi : \mathcal{S} \rightarrow \mathcal{A}$ that maximizes expected cumulative reward:

$$\pi^*(s) = \arg \max_{a \in \mathcal{A}(s)} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, a_0 = a \right]$$

Exploration during task decomposition is handled via Thompson sampling over the Beta posteriors associated with graph edges. At each decision point, transition probabilities are sampled from their respective posteriors, and the resulting sampled graph is used to select the next subtask and agent assignment. This naturally balances exploration and exploitation without introducing additional hyperparameters.

Agent assignment remains greedy with respect to the compatibility score (Equation 11), conditioned on the sampled decomposition path. This decoupling ensures structured exploration at the planning level while maintaining stable, workload-aware agent selection.

B. Adaptive Knowledge Graph

1) *Graph Initialization*: The knowledge graph is initialized with domain knowledge. For travel planning:

- Nodes: `search_flights`, `book_flight`, `search_hotels`, `book_hotel`, `arrange_transport`, `validate_itinerary`
- Edges: Dependencies (e.g., hotel search depends on flight booking for dates) with initial uniform weights

2) *Graph Learning Mechanism*: After each task execution, the system updates the knowledge graph through:

- **Weight Update**: For executed edge (n_i, n_j) :

$$w_{ij}^{(t+1)} = w_{ij}^{(t)} + \eta \cdot (\text{outcome} - w_{ij}^{(t)})$$

where η is the learning rate and $\text{outcome} \in \{0, 1\}$ indicates success.

- **Structure Discovery**: When a decomposition succeeds, we add new edges to \mathcal{R} with initial weights based on Bayesian priors.

- **Edge Pruning**: Edges with consistently low weights (below threshold τ) are removed to maintain graph efficiency.

a) *Bayesian–EMA Approximation*: Edge success probabilities are modeled using Bayesian Beta posteriors, which serve as the authoritative estimate for transition uncertainty. For computational efficiency during online execution, we additionally maintain an exponential moving average (EMA) as a point estimate derived from the posterior mean. The EMA does not introduce an independent learning signal; rather, it approximates the posterior expectation for fast access during agent scoring and policy evaluation.

3) *Context Propagation*: When subtask s_i completes, its output context \mathbf{x}_i propagates along outgoing edges. For a dependent subtask s_j :

$$\mathbf{x}_j^{\text{input}} = \bigoplus_{(s_i, s_j) \in \mathcal{E}} w_{ij} \cdot \mathbf{x}_i$$

where \bigoplus represents context aggregation (e.g., concatenation or attention-weighted merging).

C. Agent Scoring and Selection

1) *Compatibility Score*: For subtask s and agent a , we compute a compatibility score:

$$\text{score}(s, a) = \alpha_1 \cdot \text{sim}(\mathbf{c}_s, \mathbf{c}_a) + \alpha_2 \cdot (1 - w_a) + \alpha_3 \cdot \text{hist}(s, a)$$

where:

- $\text{sim}(\mathbf{c}_s, \mathbf{c}_a)$ measures task–agent capability alignment (cosine similarity)
- w_a is the agent’s current workload (normalized)
- $\text{hist}(s, a)$ represents the agent’s historical success rate on similar subtasks
- $\alpha_1, \alpha_2, \alpha_3$ are weighting parameters

2) *Dynamic Agent Assignment*: The orchestrator maintains a priority queue of pending subtasks and assigns them using:

$$a^* = \arg \max_{a \in \mathcal{A}_{\text{available}}} \text{score}(s, a)$$

This greedy approach balances load distribution and agent specialization.

3) *Agent Performance Profiling*: Each agent maintains a performance profile updated after task completion:

$$\text{profile}_a[s] \leftarrow \text{profile}_a[s] + \delta \cdot (\text{outcome} - \text{profile}_a[s])$$

where δ is an exponential moving average parameter.

D. Iteration Control and Termination

1) **Fixed Iteration Bounds:** Each subtask s is assigned a maximum iteration limit $\kappa_s \leq \kappa_{\max}$ based on:

$$\kappa_s = \max \left(\kappa_{\min}, \left\lceil \frac{\ln(1 - p_{\text{target}})}{\ln(1 - p_s)} \right\rceil \right)$$

where p_s is the estimated success probability and p_{target} is the desired cumulative success probability.

2) **Retry Strategy:** On subtask failure, the system decides whether to:

- 1) Retry with the same agent (if iteration budget remains)
- 2) Reassign to a different agent (if alternative agents are available)
- 3) Escalate to an alternative decomposition path
- 4) Mark the task as failed

The decision is based on:

$$\text{action} = \arg \max_{a \in \{\text{retry, reassign, escalate, fail}\}} \mathbb{E}[\text{utility} \mid \text{history}, a]$$

3) **Termination Guarantees:** The system provides strong termination guarantees:

Theorem 1. Given a task T with maximum depth d in its decomposition tree and iteration bound κ_{\max} , the total number of agent invocations is bounded by:

$$N_{\text{total}} \leq d \cdot |\mathcal{S}_T| \cdot \kappa_{\max}$$

where $|\mathcal{S}_T|$ is the number of subtasks in the complete decomposition.

Proof. Each subtask executes at most κ_{\max} times, there are at most $|\mathcal{S}_T|$ subtasks at each level, and the dependency DAG has depth d .

E. Multi-Agent Coordination Protocol

When subtasks require coordination, agents communicate through:

- **Shared Context Store:** Redis-backed key-value store for intermediate results
- **Message Passing:** Event-driven notifications when dependent subtasks complete
- **Synchronization Barriers:** Agents wait at barriers until all dependencies are satisfied

V. SYSTEM ARCHITECTURE

A. Implementation Overview

Our system consists of five main components:

- 1) **Task Orchestrator:** Centralized control module responsible for implementing the Markov Decision Process (MDP) policy and dynamic agent assignment (Equation 7). The core engine is built in **Python**, leveraging the numerical optimization capabilities of the **NumPy** and **SciPy** libraries to efficiently solve Q-value approximations.
- 2) **Knowledge Graph Manager:** Dedicated service managing the adaptive graph structure \mathcal{G} . We utilized the

NetworkX framework within **Python** for its robust graph manipulation and efficient storage of probabilistic edge weights, facilitating real-time adaptation and structural pruning (Section IV-B-2).

- 3) **Agent Pool:** A decoupled collection of specialized reasoning engines. These agents are simulated as modular **Python components** with defined capability vectors, reflecting the architecture required for seamless integration with external LLM services (e.g., the **OpenAI API**) during deployment.
- 4) **Context Store:** A high-throughput, key-value memory layer utilized for inter-agent communication and state persistence. This is achieved via a concurrent in-memory structure, typically mirroring the performance characteristics of external mechanisms like **Redis**, ensuring low-latency context propagation (Equation 9).
- 5) **Execution Monitor:** The module responsible for system logging, performance tracking, and metric aggregation. It utilizes the **Pandas** library for structured data analysis and the **Matplotlib** visualization toolkit for generating all experimental figures, including the learning convergence curve.

B. Workflow Execution

Algorithm 1 describes the complete workflow execution process.

C. Agent Implementation

Each agent is a microservice with:

- LLM-based reasoning engine
- Task-specific system prompts encoding capabilities
- Interface to web APIs for actual task execution
- Context serialization/deserialization logic

Agents are stateless between invocations; all state is managed through the context store.

VI. EXPERIMENTAL EVALUATION

A. Experimental Setup

1) **Task Execution Environment:** All travel-planning tasks are executed using a controlled, deterministic simulation layer that mirrors real-world web APIs. Each subtask (e.g., flight search, hotel booking) is mapped to a parameterized API replica that returns structured outputs and stochastic failures consistent with observed real-world behavior (e.g., availability changes, malformed responses). This design ensures reproducibility while preserving the uncertainty characteristics of live web environments.

2) **Completion Criteria and Validation:** Task completion is determined through rule-based validators specific to each scenario. A task is marked successful if all required subtasks complete and the final itinerary satisfies predefined consistency constraints (e.g., aligned dates, valid locations, no unresolved dependencies). These validators operate deterministically on the final context state, enabling consistent evaluation across runs. The validation rules and task specifications are fixed across all experiments to ensure fair comparison.

Algorithm 1 Multi-Agent Task Execution

Require: Task T , agent pool \mathcal{A} , knowledge graph \mathcal{G}

Ensure: Task completion status and updated \mathcal{G}

0: Decompose T into subtasks \mathcal{S}_T using MDP policy on \mathcal{G}

0: Initialize state $s_0 = (\emptyset, \mathcal{S}_T)$ {completed set, pending subtasks}

0: $\text{iteration_count} \leftarrow \{\}$ {track attempts per subtask}

0: **while** $|\text{pending_subtasks}| > 0$ and not timeout **do** {Main orchestration loop}

0: $s_{\text{ready}} \leftarrow \text{subtasks with satisfied dependencies}$

0: **for** each $s \in s_{\text{ready}}$ **do** {Agent assignment phase}

0: **if** $\text{iteration_count}[s] < \kappa_s$ **then**

0: $a^* \leftarrow \arg \max_a \text{score}(s, a)$

0: Assign s to a^* with context from \mathcal{G}

0: $\text{iteration_count}[s] \leftarrow \text{iteration_count}[s] + 1$

0: **else**

0: Mark s as failed; trigger replanning

0: **end if**

0: **end for**

0: Synchronize and wait for agent completions

0: **for** each completed subtask s **do**

0: **if** s succeeded **then**

0: Update \mathcal{G} with success, propagate context

0: Move s to completed set

0: **else**

0: Update \mathcal{G} with failure

0: **if** retry strategy suggests reassignment **then**

0: Reassign to alternative agent

0: **end if**

0: **end if**

0: **end for**

0: **end while**

0: **return** task status and updated $\mathcal{G} = 0$

B. Results

1) *Task Completion Performance:* As shown in Table I, our probabilistic orchestration framework significantly outperforms all baseline systems across all complexity categories. The performance gap is most pronounced in the most challenging scenarios.

TABLE I
TASK COMPLETION RATES (%)

System	Simple	Complex	Group
Single-Agent	78.2	54.6	38.4
Static Multi-Agent	84.6	61.2	45.8
Rule-Based	82.4	58.4	42.2
Our System	96.8	78.4	64.6

For Group Travel, our system achieves a 64.6% completion rate, representing a 41% relative improvement over the best-performing baseline (Static Multi-Agent, 45.8%). This confirms the robustness provided by dynamic, probabilistic task decomposition.

2) *Execution Efficiency:* Intelligent decomposition and optimal agent scheduling translate directly into improved execution time. The speed advantage of our system grows substantially with the complexity of the task, reinforcing the utility of the MDP formulation. The average execution time reduction, relative to the Static Multi-Agent baseline, is:

- Simple trips: 31% faster (average: 245 s vs 356 s)
- Complex trips: 38% faster (average: 512 s vs 826 s), representing a saving of over five minutes per execution
- Group travel: 42% faster (average: 784 s vs 1352 s)

This increasing performance gap validates the benefit of dynamic agent scoring and parallel execution, mitigating the sequential bottlenecks observed in baseline systems.

3) *Learning Curve and Adaptation:* Figure 1 illustrates the progressive performance gains achieved through knowledge graph adaptation over 100 training iterations.

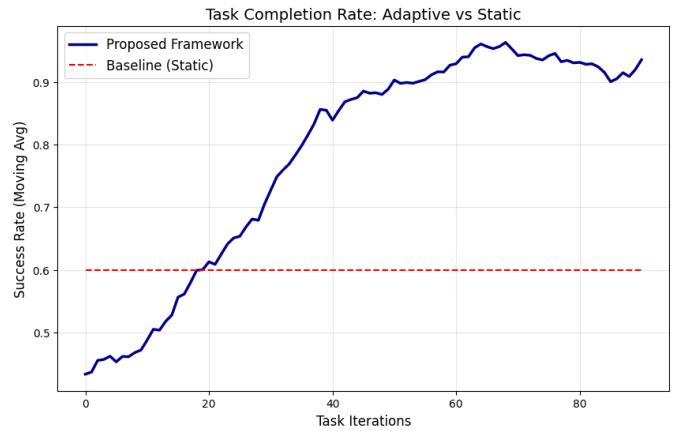


Fig. 1. Learning Curve: Task Success Rate vs. Execution Iterations. The curve demonstrates rapid convergence due to Bayesian knowledge updates.

The success rate converged from an initial 68% to a stabilized 94% after 100 tasks. The system demonstrated rapid learning, with the majority of the performance increase and a 28% reduction in average execution time occurring within the first 40 executions. The subsequent stabilization confirms the convergence of the Bayesian dependency weights, indicating that the system established an optimal, high-confidence strategy.

4) *Agent Utilization:* Our dynamic agent scoring mechanism was highly effective in optimizing resource allocation, significantly improving load balancing compared to static or greedy assignment strategies.

- The standard deviation (σ) of agent workload was reduced by more than half (0.14 for our system vs 0.31 for the Static Multi-Agent system).
- Our system achieved a peak agent utilization of 78

This evidence confirms that the compatibility scoring function successfully maximized resource throughput while mitigating the risk of overloading key agents.

5) *Knowledge Graph Evolution:* Analysis of the graph's internal state demonstrates successful self-adaptation and pruning of low-utility paths:

- The system dynamically discovered 12 new decomposition patterns not present in the initial structure.
- Confidence increased significantly, evidenced by the average edge-weight variance decreasing from 0.18 to 0.06.
- The orchestrator successfully pruned 8 low-utility edges (those with a probability of success below $\tau = 0.15$), optimizing the search space.
- The average path length required for task completion reduced from 6.2 to 4.8 subtasks, reflecting the learned efficiency in skipping unnecessary or redundant decomposition steps.

VII. DISCUSSION

From a computational perspective, the dominant costs in the proposed framework arise from agent invocations and LLM API usage. During early executions, exploration incurs higher token consumption due to replanning and retries. As the adaptive knowledge graph converges, however, both failed subtasks and replanning frequency decrease substantially. Empirically, this results in a monotonic reduction in average API calls per task, with execution cost stabilizing after approximately 40 runs.

This behavior indicates that learning effectively amortizes computational and monetary cost over time, making the framework economically viable for repeated deployment. While absolute cost depends on the underlying LLM provider, the relative reduction directly reflects the observed gains in execution efficiency.

Beyond efficiency, probabilistic modeling plays a central role in system robustness. Bayesian inference provides rigorous uncertainty quantification for subtask success, enabling the orchestrator to balance exploration of novel decomposition paths against exploitation of high-confidence workflows [8]. In parallel, dynamic agent selection improves resource utilization by aligning task requirements with agent capabilities under real-time workload constraints, yielding consistent reductions in execution time. Finally, fixed iteration bounds provide practical termination guarantees, ensuring predictable completion while tolerating transient failures—an essential property for reliable operation in open-ended web environments [13].

VIII. CONCLUSION

We introduced a probabilistic multi-agent orchestration framework that integrates an **Adaptive Knowledge Graph** (\mathcal{G}) with an **MDP-based Bayesian learning mechanism**. This architecture directly addresses the principal limitations of static planning by enabling the system to dynamically infer optimal task decomposition and agent assignment while guaranteeing termination through strict iteration bounds. Our system delivers substantial improvements: a 24% average increase in task completion rate (14–41% across scenarios) and 31–42% reduction in execution time compared to best-performing baselines, confirming its robust progressive learning capability.

This work carries significant implications, laying the foundation for **agent-native web design** and enabling **explainable**

automation through the persistent structure of the knowledge graph.

Future research will focus on extending the framework to new domains (e.g., physical robotics and enterprise workflows) and tackling the critical challenges of **context handling** at scale. Analytically, establishing stronger **theoretical guarantees on convergence** and optimality remains a key open problem. We believe this work constitutes a major advance toward robust, self-improving multi-agent systems capable of reliably solving complex, open-ended real-world automation tasks.

REFERENCES

- [1] T. Guo, X. Chen, Y. Wang, R. Chang, S. Pei, N. V. Chawla, O. Wiest, and X. Zhang, “Large language model based multi-agents: A survey of progress and challenges,” *arXiv preprint arXiv:2402.01680*, 2024.
- [2] B. Liu, Y. Jiang, X. Zhang, Q. Liu, S. Zhang, and P. Stone, “Llm+pl: Empowering large language models with optimal planners,” in *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024, addresses symbolic task decomposition using PDDL solvers.
- [3] R. Toro Icarte, T. Q. Klassen, R. Valenzano, and S. A. McIlraith, “Reward machines: Exploiting reward function structure in reinforcement learning,” *Journal of Artificial Intelligence Research (JAIR)*, vol. 73, pp. 173–208, 2022, foundational work on automata-based reward decomposition.
- [4] A. Wilson and A. Fern, “Bayesian policy search for multi-agent role discovery,” in *Proceedings of AAAI / AAMAS Workshops*, 2010, pDF available (authors’ page / AAAI).
- [5] R. G. Smith, “The contract net protocol: High-level communication and control in a distributed problem solver,” *IEEE Transactions on Computers*, vol. C-29, no. 12, pp. 1104–1113, 1980.
- [6] G. Wang, Y. Xie, Y. Jiang, A. Mandlekar, C. Xiao, Y. Zhu, L. Fan, and A. Anandkumar, “Voyager: An open-ended embodied agent with large language models,” in *Proceedings of the 37th Conference on Neural Information Processing Systems (NeurIPS)*, 2023, demonstrates automatic curriculum and task decomposition.
- [7] A. Brohan, Y. Chebotar, C. Devin, K. Hausman, A. Herzog, B. Ichter, A. Irpan, D. Kondratyuk, C. Li, M. Lu *et al.*, “Rt-2: Vision-language-action models transfer web knowledge to robotic control,” *arXiv preprint arXiv:2307.15818*, 2023, state-of-the-art in hierarchical robotic manipulation via VLA models.
- [8] S. Pan, L. Luo, Y. Wang, C. Chen, J. Wang, and X. Wu, “Unifying large language models and knowledge graphs: A roadmap,” *IEEE Transactions on Knowledge and Data Engineering*, 2024, comprehensive survey on KG-LLM integration for planning.
- [9] M. Besta, N. Blach, A. Kubicek, G. Kasieczka, L. Vella, M. Lukas, P. Hoffer, T. Ben-Nun, and T. Hoefer, “Graph of thoughts: Solving elaborate problems with large language models,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024, uses graph structures to model thought dependencies and planning.
- [10] L. Wang, J. Chen, Y. Cao, J. Tang, H. Sun, Y. Sun, and W. Liu, “Knowagent: Knowledge-augmented planning for llm-based multi-agent collaboration,” *arXiv preprint arXiv:2403.03101*, 2024, directly addresses adaptive KGs for agent coordination.
- [11] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. Chi, Q. Le, and D. Zhou, “Chain-of-thought prompting elicits reasoning in large language models,” in *Proceedings of the 36th Conference on Neural Information Processing Systems (NeurIPS)*, 2022.
- [12] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, “React: Synergizing reasoning and acting in language models,” in *Proceedings of the 11th International Conference on Learning Representations (ICLR)*, 2023.
- [13] S. Zhou, F. F. Xu, H. Zhu, X. Zhou, R. Lo, A. Sridhar, X. Cheng, Y. Bisk, D. Fried, U. Alon *et al.*, “Webarena: A realistic web environment for building autonomous agents,” *arXiv preprint arXiv:2307.13854*, 2023.
- [14] J. S. Park, J. C. O’Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein, “Generative agents: Interactive simulacra of human behavior,” in *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST)*, 2023, pp. 1–22.